

Operating System

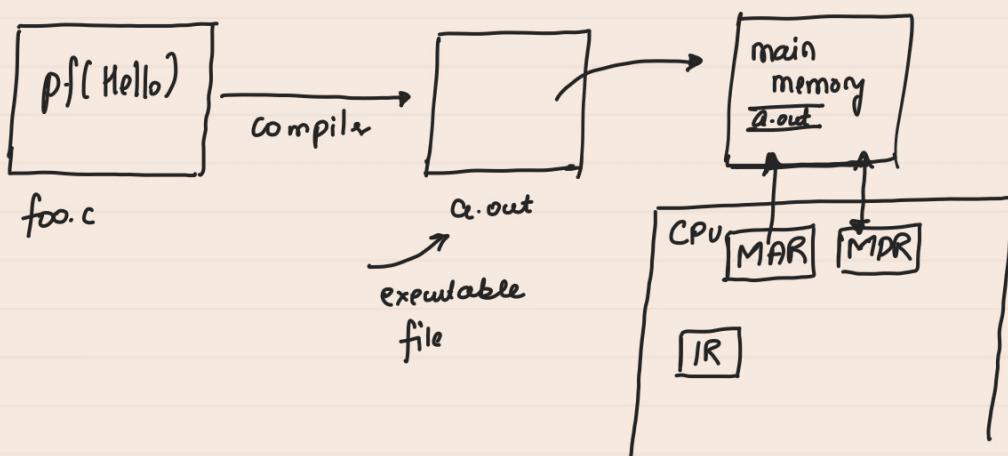
Syllabus:

- System calls, processes, threads, inter-process communication, concurrency and synchronization
- deadlock
- CPU and I/O scheduling
- Memory management and virtual memory
- file systems

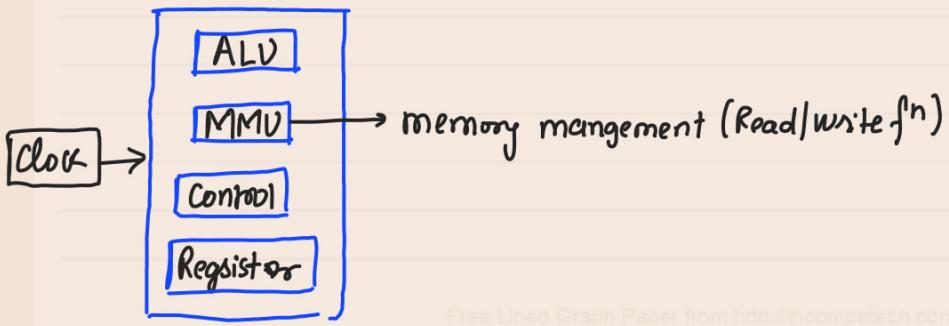
lecture 1

Prerequisite: The journey of program (50Mins)

COA and OS go together



Inside CPU/processor



Free Lined Graph Paper from <http://incompetech.com/graphpaper/lined/>

CPU fetch 7 while is loop till the program (set of instruction gets completed)

- Q. What happens in CPU?
- fetch
 - decode
 - execute
- will go in loop till the program (or) instructions get completed.
- infinite cycle

Q. Once you compile, you can run the file ??

Ans: Whenever, you want to run.

running simple program in CPU

```
int a=20, b=30;
int c;
c = a+b;
```

→ PC increment karna hoga manually karna hoga
 → CPU will fetch, load, execute
 $c \leftarrow a+b$; (ALU)

Q. Why "printf" needs OS support??

- because CPU doesn't own the monitor
- monitor is external (called I/O device)
- CPU has to take permission from monitors

Q. Why need OS?

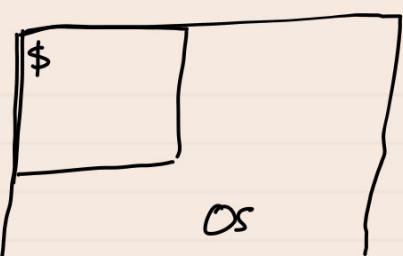
- manually too hectic
- these tasks are very common task
- even if we ask user to do it manually then first they need to be computer scientist XD.

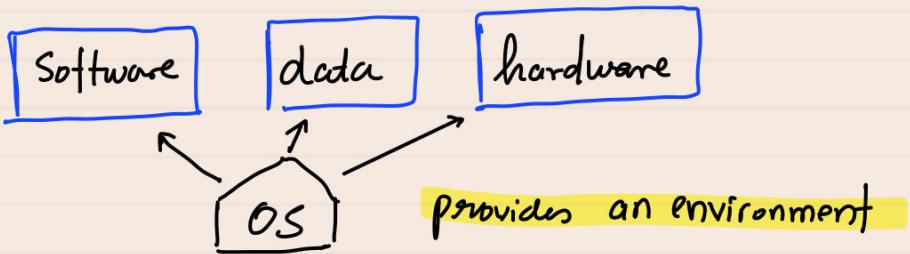
Q. What OS does?

- manage memory
- manage resources (I/O devices)
- manage file system

What is shell in OS

Lined Graph Paper from <http://incompetech.com/graphpaper/lined/>





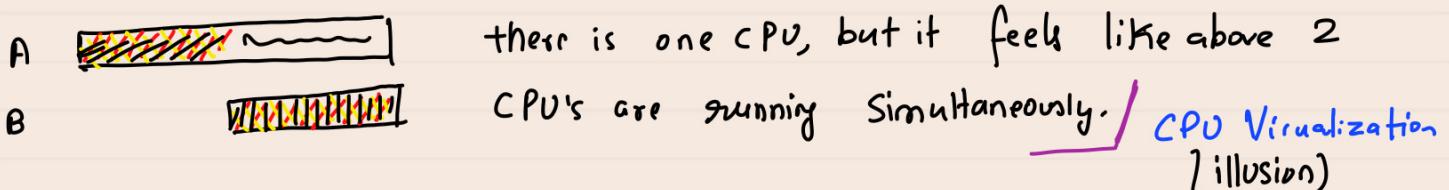
$4 \text{ core} \equiv 4 \text{ CPU}$

DS - Provides (Virtualization)

lecture 2

- Q. What happens when you switch on the computer?
- Q. How does even OS code loaded into memory?
- Q. How to create a new process
↳ how OS load the process in the main memory?
- Q. Is it possible to tweak DS in our personal laptop?

Ans Yes possible



HOW TO CREATE NEW PROCESS

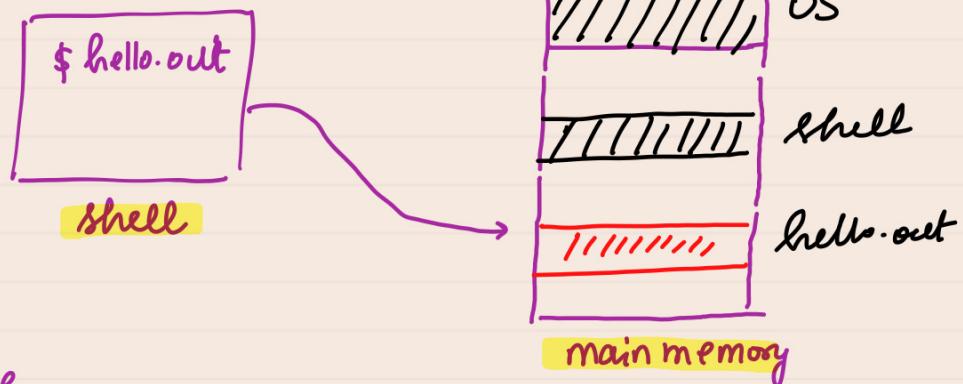
- ① User writes a program
- ② Compiler converts the program into executable code (byte code)
- ③ User requires the OS to run the byte code.
- ④ OS "load" the program into memory
- ⑤ OS sets registers properly and starts running the code

- DS is always in main memory (loaded by bootloader at start of computer), but it does not mean that DS is running all the time.

example : `printf(" ")`
→ this line can't be executed without the interference of OS!

malloc()
OS needed!

fork-exec()



Objective

Understand how these two steps are performed?

- ① Create Space for hello.out
- ② Load hello.out in that space.

Free Lined Graph Paper from <http://incompetech.com/graphpaper/lined/>

Process

- a program in execution
- processes compete for resources
- context of process
- Multiple processes at same time

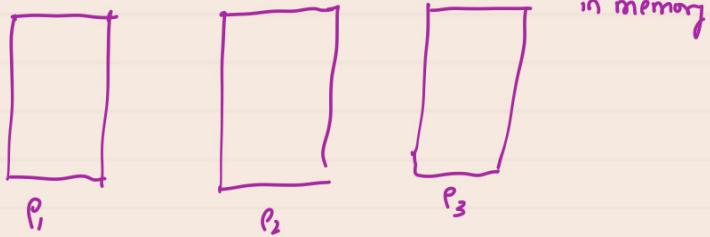
Context of process (pcb)

- Every process has context and OS keeps track of that.
- like a school keep tracks of activities of there students, an OS keep

track of its processes.

What all things does OS track about process??

- ① PC value
- ② general purpose register
- ③ file opened etc



struct pcb {

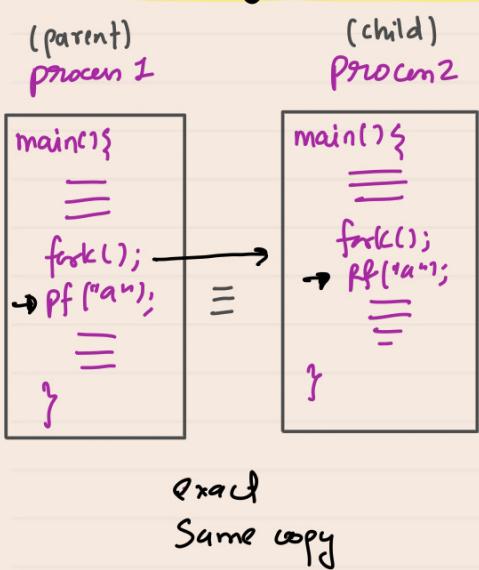
```

  _____
  |
  |
  |
  |
  |
  }

```

pcb: a data structure that keep the track of process.

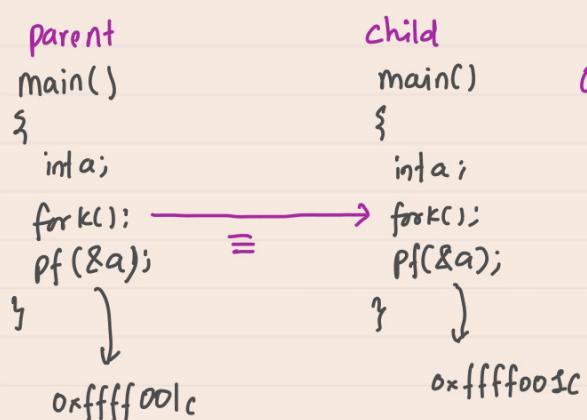
fork() : System Call



fork():

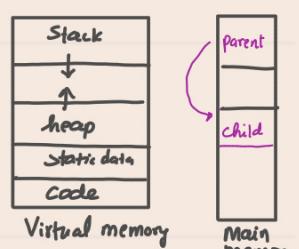
- ① create EXACT same copy
- ② Parent and child will start executing after the line which had fork()
- ③ fork() will return some integer value to both process
- ④ fork return "0" to child
- ⑤ fork return nonzero to parent
- ⑥ fork create exact same copy hence logical address same but physically variables are stored at different locations.

o Logical address and actual address (postponed)



Q: Will there address be same?

Ans: Will print exactly same thing!!



M: (Read My Question): Why on earth are we doing so? What are getting by doing this, aren't we using resources wastefully?

```

#include <stdio.h>
#include <unistd.h>

int main() {
    int pid = fork();

    if(pid == 0){
        printf("I am child\n");
    }
    else{
        printf("I am parent\n");
    }
    return 0;
}

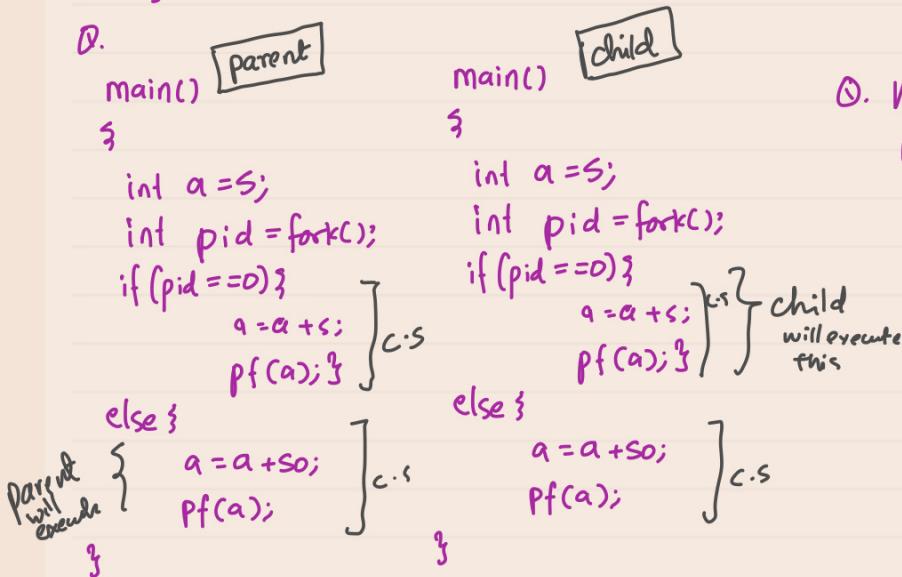
```

OSCodes - zsh - 80x24
(base) sachinmittal@Sachins-MacBook-Pro OSCodes % gcc
(base) sachinmittal@Sachins-MacBook-Pro OSCodes % ./a
I am parent
I am child
(base) sachinmittal@Sachins-MacBook-Pro OSCodes %

Q. How many times each of them will be printed?

```
int main()
{
    printf("GATE\n");
    — 1 time (i.g)
    int pid=fork();
    if(pid==0)
        printf("GATEOverflow!");
        — 1 time
    }
    else
        printf("Go Class\n");
    }
    printf("IISCE\n");
    ——————> 2 time
}
```

Q.



{ how many patterns / combination we will discuss later }

Q. What will child print and what will parent print ??

(10, 60)
↑
Child parent
(60, 55)
↑
child parent

Q. Child → parent, order

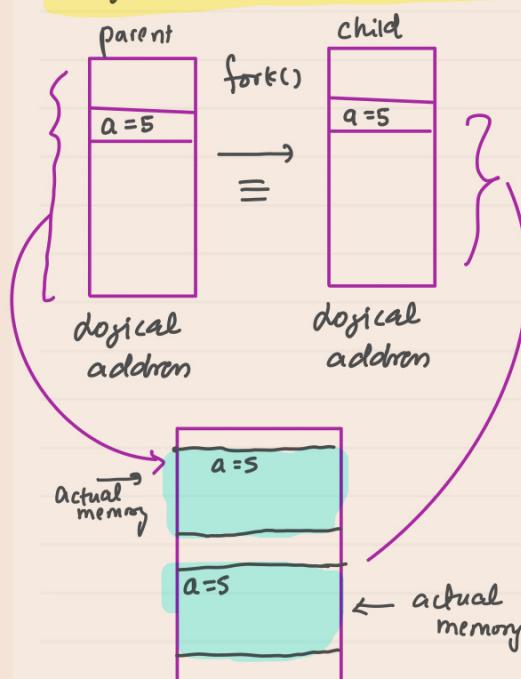
$a=5$ $a=5$
 $a_c=10$ or $a_c=10$
 $a_p=55$ $a_p=60$ (taken $a=10$)

(✓) (✗)

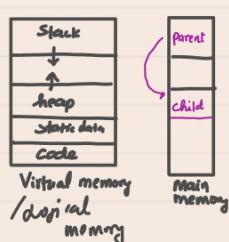
Free Lined Graph Paper from <http://tiny.cc/meyarw> | graphpaperonline

"It doesn't matter what you are doing to child or parent execution order doesn't matter"
Why? how?

• Logical address and actual address (postponed)



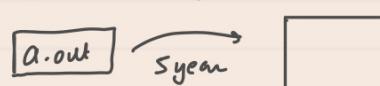
"Exactly same logical memory but different actual memory"



MOV R1,

1000

• Why "logical" memory (postponed) (logical memory)



```
Consider the following code fragment:
if (fork() == 0)
{
    a = a + 5;
    printf("%d, %p\n", a, &a);
}
else
{
    a = a - 5;
    printf ("%d, %p\n", a,&a);
}
```

u,v - parent process
x,y - child process

u=a+5
x=a-5
u-5=x+5;
u=x+10
x=u+5
x=u+10
(✓)

child parent

same, v=y

Let u, v be the values printed by the parent process and x, y be the values printed by the child process. Which one of the following is TRUE?

- A. $u = x + 10$ and $v = y$
- B. $u = x + 10$ and $v = y$
- C. $u + 10 = x$ and $v = y$ (correct)
- D. $u + 10 = x$ and $v = y$

even if global variable

Do you think
Compiler will think

Where in main memory
Space will be free??
NOPE. It logically
think there is complete
continuous memory free

main memory

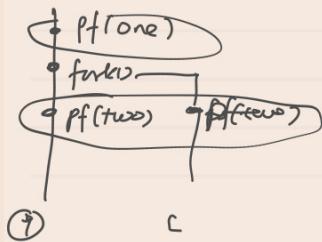


address space / logical memory

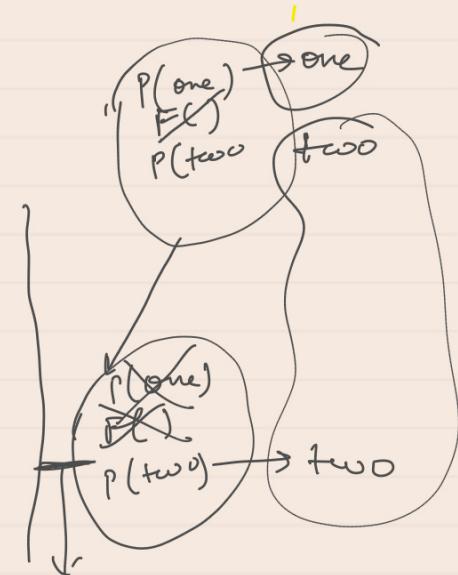
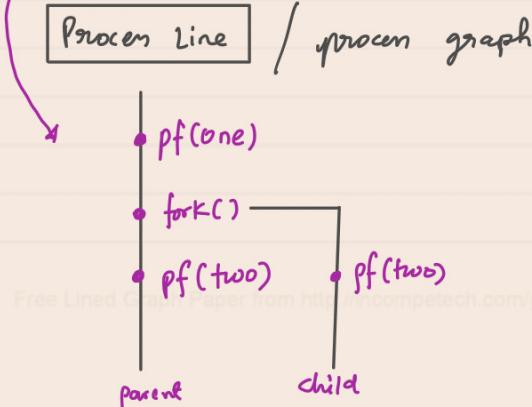
Q. What will be output of the following program?

```
int main()
```

```
printf("one\n");
fork();
printf("two\n");
return 0;
```



Ans: one
two ← parent | child (order!)
two ← child | parent



Free Lined Graph Paper from <http://www.competech.com/graphpaper/lined/>

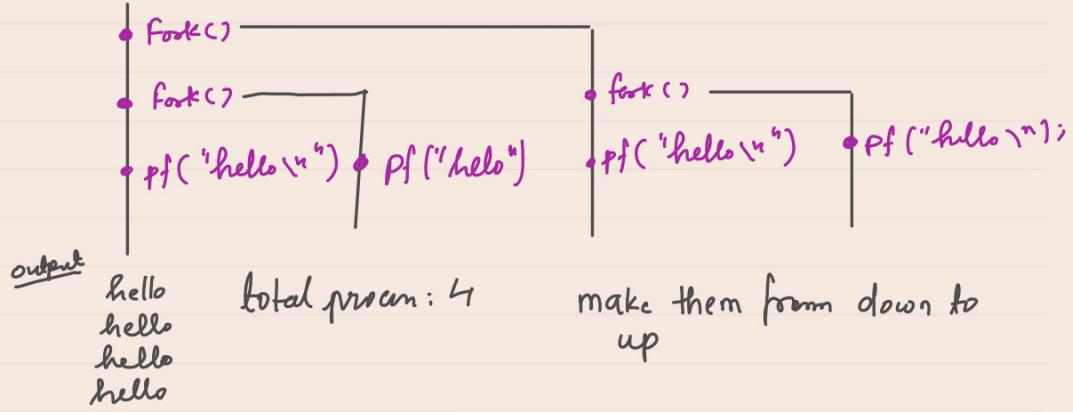
Q. int main()

```
Fork();
Fork();
printf("Hello\n");
exit(0);
```

}

output??

output process line / process graph



* Q. how many Hello's printed. (Assume fork does not fail.)

```
int main()
```

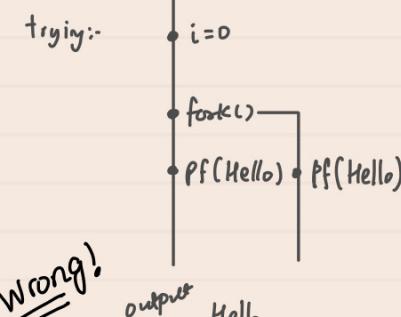
```
int i;
for(i=0; i<2; i++) {
    fork();
    printf ("Hello\n");
}
```

}

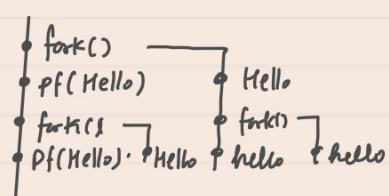
return 0;

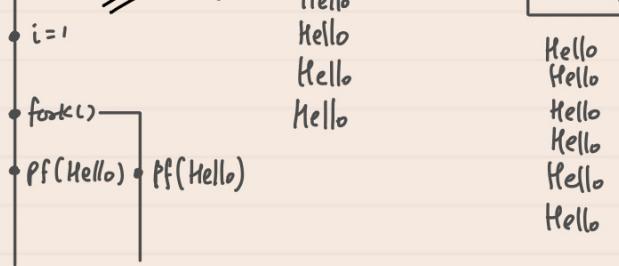
- A. 2
- B. 4
- C. 6
- D. 8

trying:-



sir's method





Process: 6

Q. How many 'Hello' printed? (Assume fork() do not fail)

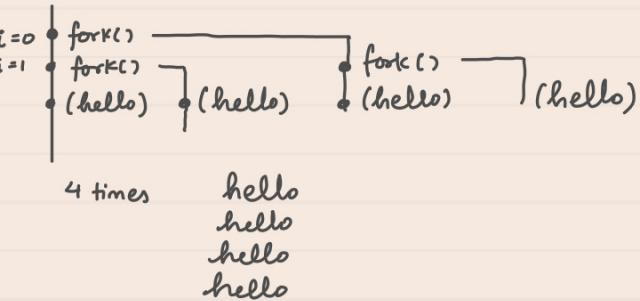
int main()

```

int i;
for(i=0; i<2; i++) {
    fork();
    printf("Hello\n");
}
return 0;

```

- a. 2
- b. 4
- c. 6
- d. 8

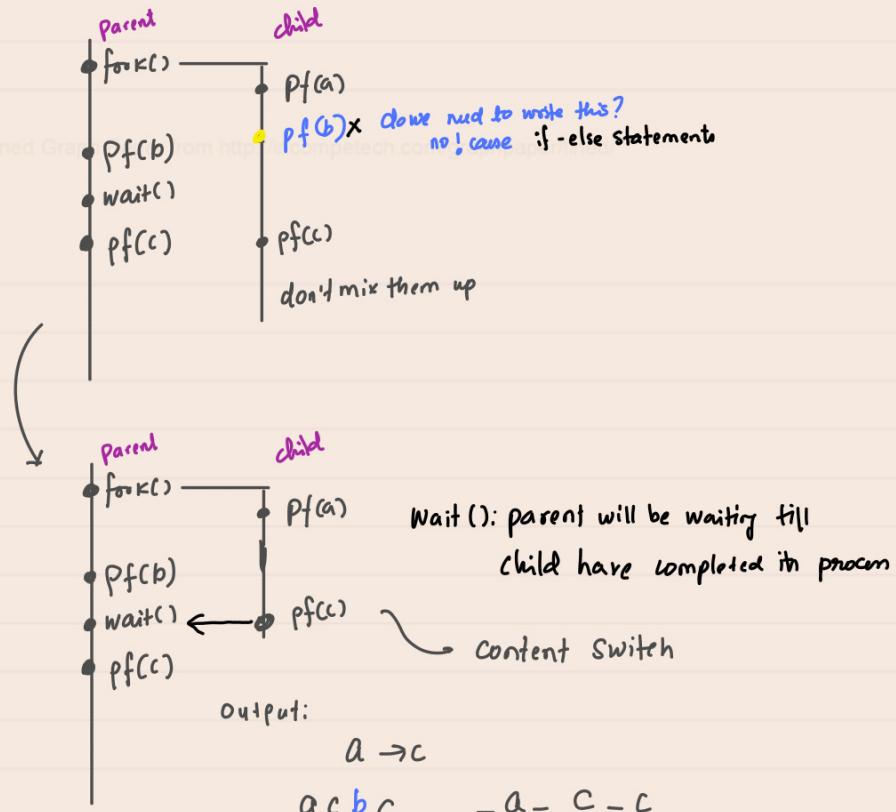


* Q4 Write down all the possible patterns printed by following code

```

int main() {
    if (fork() == 0) { // child
        printf("a");
    } else { // parent
        printf("b");
        wait();
    }
    printf("c");
    exit(0);
}

```



Q.8 How many processes are created with these fork()

Statements?

HOMEWORK (Wrong!)

```

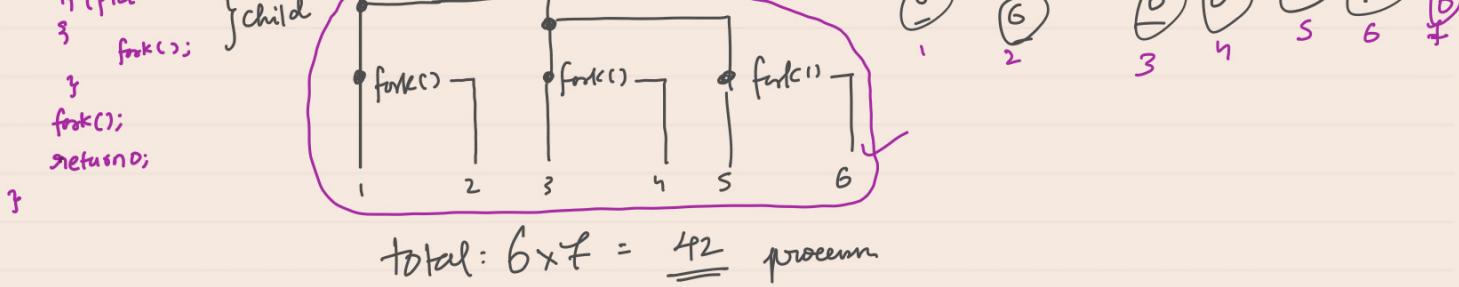
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void) {
    pid_t pid;
    pid = fork();
    pid = fork();
    pid = fork();
    if (pid == 0)

```

parent pid-f pid = fork()





Lecture 3

Q. How many child processes will create?

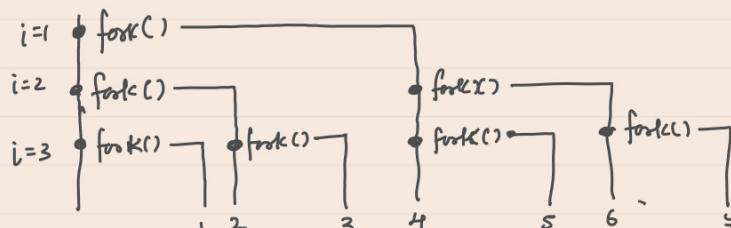
```
for(i=1; i<=3; i++)  

{  

  fork();  

}
```

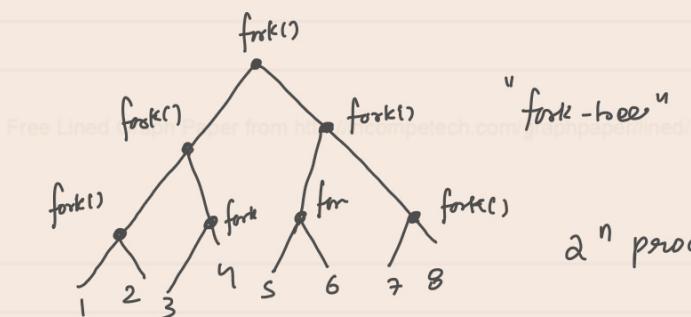
Try it



Ans: 4 child processes

1 parent process

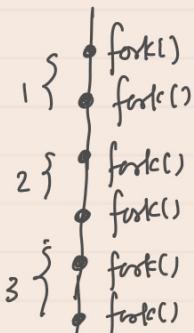
Total process: 8



Q. How many child process will create?

```
for(i=1; i<=3; i++) {
```

 fork();
 fork();
 }



$$2^6 = 2^2 \times 2^2 \times 2^2$$

$$4 \times 4 \times 4$$

$$16 \times 4 = 64$$

No. of child process: 63

Q. $\text{for}(i=0; i < n; i++) \text{ fork();}$

Total no. of child processes?

"n times" $\rightarrow 2^n$ processes

$2^n - 1$ child processes

Q.8 How many processes are created with these fork()

Statements?

HOMWORK (Wrong!)

```
#include<stdio.h>
```

```
#include<stdlib.h>
```

```
#include<unistd.h>
```

```
int main(void){
```

```
  pid_t pid = fork();
```

```
  pid = fork();
```

```
  pid = fork();
```

* * *

*

*

*

*

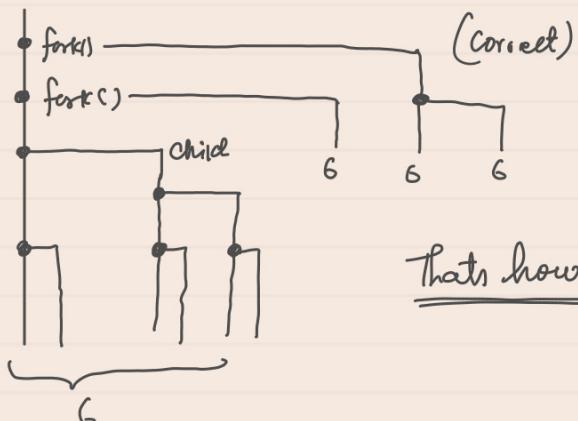
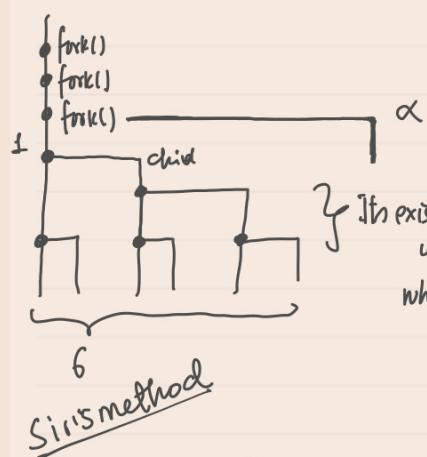
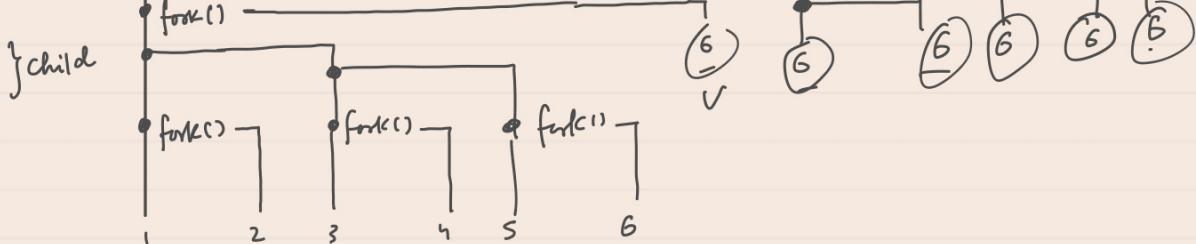
parent pid-f pid = fork()



```

pid = fork();
if (pid == 0)
{
    fork();
}
fork();
return 0;
}

```



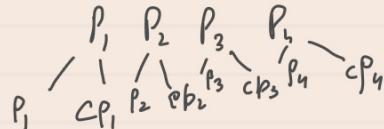
```

int main(void){
    pid_t pid = fork();
    pid = fork();
    pid = fork();
    if (pid == 0)
    {
        fork();
    }
    fork();
    return 0;
}

```

how many processes? = 4

how many processes? = 8



$P_1, P_2, P_3, P_4, c(P_1, P_2, P_3, P_4) \rightarrow \text{zero pid}$

non-zero pid



$$4 \times 2 = 8$$



$$2^2 = 4$$

$$4 \times 4 = 16$$

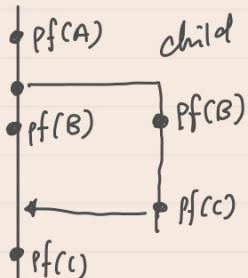
$$\therefore \text{Total} = 8 + 16 = 24$$

Q. Consider the following example program. List all legal outputs this program may produce when executed on a unix system. The output consists of strings made up of multiple letters.

```

#include <unistd.h>
#include <sys/wait.h>
// W(A) means write(1, "A", sizeof("A"))
#define W(x) write(1, #x, sizeof(#x))
int main(){
    pf(A);
    int child = fork();
    W(B); — pf(B)
    if (child)
        wait(NULL);
    W(C);
}

```



A _ B - C - C

A B B C C

A B B C c

A C B C :

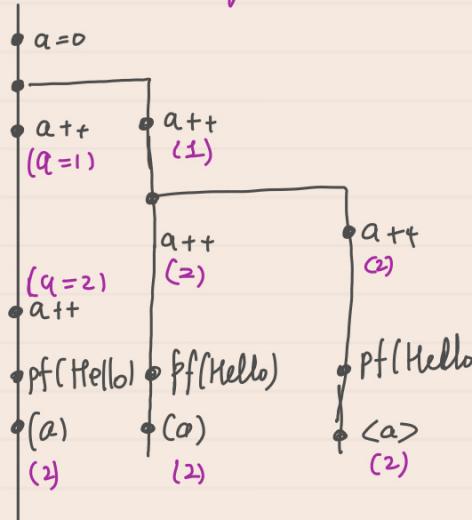
∴ total 2 patterns!

if (non-zero) = TRUE

if (zero) = False

Q. how many times does the following program print "Hello!"?
also value of a?

```
int a=0;
int rc = fork();
a++;
if (rc==0) {
    rc = fork();
    a++;
}
else {
    a++;
}
printf("Hello!\n");
printf("a is %d\n", a);
```



Free Lined Graph Paper: <http://www.comgraphpaper.com/graphpaper/lined/>

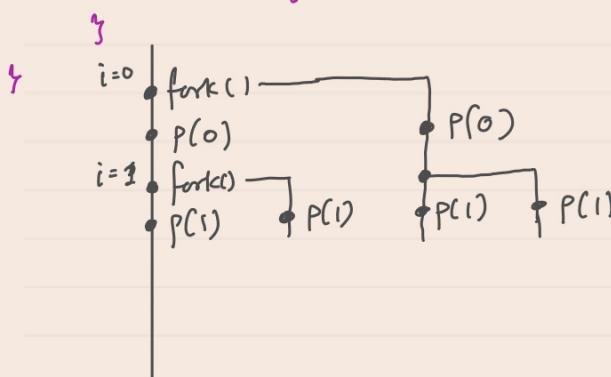
Ans 3 times "Hello"?

2 max value?

Q. int main()

```
int i=0;
while (i<2) {
    fork();
    printf("%d", i);
    i=i+1;
}
```

X A: 0 1 0 1
✓ B: 0 0 1 1 1 1
✓ C: 0 1 0 1 1 1
X D: 0 1 1 1 0 1
E: 0 1 1 0 1 1 ↳ need 0 to access 1's



Q. int main()

(H.W.)

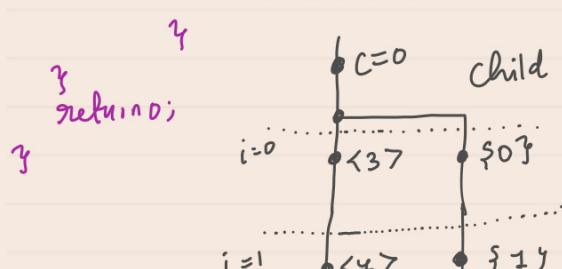
```
int c=0;
for (int i=0; i<3; i++) {
    if (fork() == 0) {
        printf("%d", c+i);
        fflush(stdout);
    }
    else {
        c=c+i+3;
    }
}
```

a) how many times fork() involved?
b) correct Outputs

(Later)

- a. 0 1 4 6 5 9 2
- b. 0 1 5 4 9 6 2
- c. 0 4 2 6 1 5 2
- d. 9 4 0 5 1 2 6
- e. 4 0 1 6 9 2 5
- f. 6 4 9 5 0 1 2

attempted but unable to understand



c = 2

EXEC SYSTEM CALL

exec() → parameter → any **executable** file

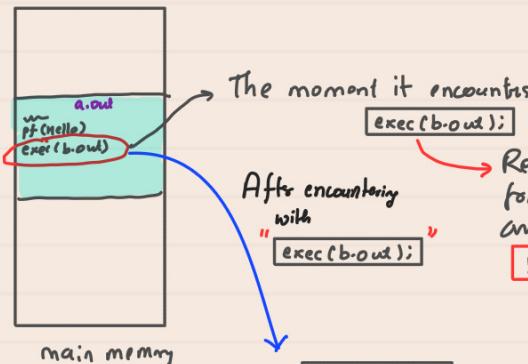
Free Lined Graph Paper from <http://incompetech.com/graphpaper/lined/>

Suppose

```
a.out
int main() {
    ...
    pf("Hi Adib");
    ...
    exec(b.out);
    ...
    pf("Adib failed");
    ...
}
```

b.out

```
int main() {
    ...
    pf("You got this!");
    ...
}
```



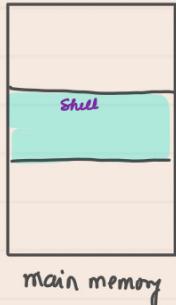
What is the output if we ran
a.out file ??

Hi Adib	Hi Adib
You got this!	You got this!
(✓)	(X)
Adib failed → never printed!!	

a.out removed from
main memory!



Process Creation



① Suppose, **shell** is running in main memory, and in cmd you entered **./a.out**, then what?

shell code (somewhat)

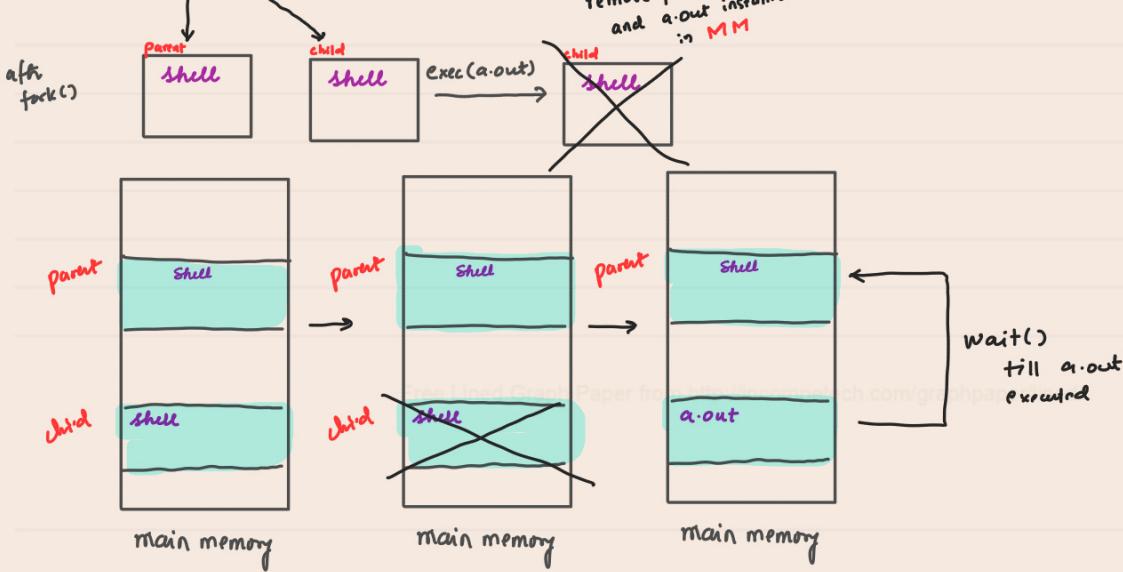
```
while (TRUE) {
    Some command
    scanf("%s", &t);
    pid = fork();
    if (pid == 0) // child
        exec(t);
    else
        wait();
}
```

(Not the only, UNIX does
this way, other OS may
do other way too!)

(discussing Idea!)



remove from MM



You write a UNIX shell, but instead of calling fork() then exec() to launch a new job, you instead insert a subtle difference: the code first calls exec() and then calls fork() like the following:

```
shell (...) {
    ...
    exec (cmd, args);
    fork();
    ...
}
```

Does it work? What is the impact of this change to the shell, if any? (Explain)

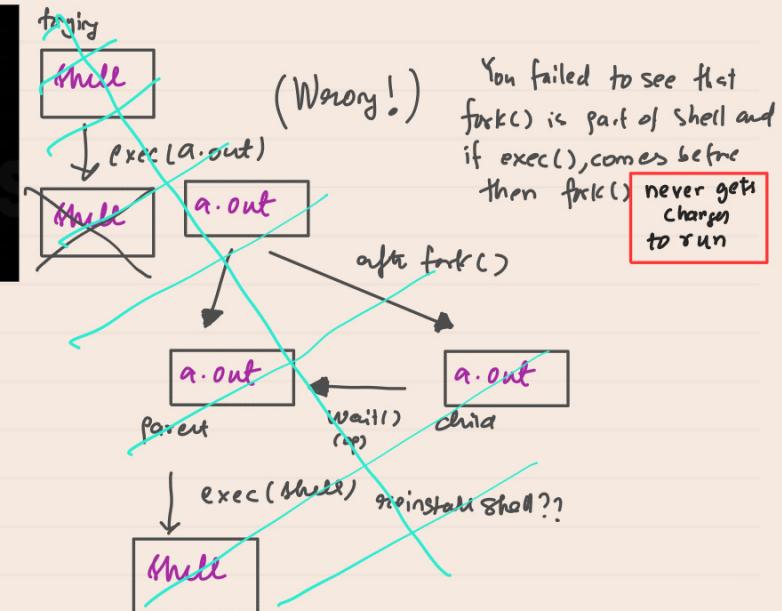
→ Means Working just like "earlier"??



Who creates shell
Inside main memory??

init (Bootload kernel)
↳ Shell

[process tree of unix]
↳ data



(Correct)

Ans: No.

- We are destroying the shell.
- We are just replacing shell with a.out
- fork() never gets executed!!

Part (a) [3 MARKS] whatA

```
int main() {
    int i = 0;
    printf("%d ", i);
    i = i + 1;
    if (i >= 2)
        exit(0);
    execv("./whatA");
}

A: 0 1 2
B: 0 0 0
C: 0 0 0 0 0 ... (forever) (✓)
D: 0 1 2 0 1 2 ... (forever)
E: none of the above
```

initially "0" printed, $i=1$ (updated)
as $i \geq 2$ so goes to next line exec(whatA),
which means that it will load whatA in place
of whatA, and again $i=0$, "0" printed

∴ 0 0 0 0 ... (forever)

- > What is System call?
> why they are always in limelight?

System Calls

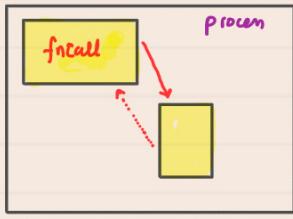
function call

VS

System call

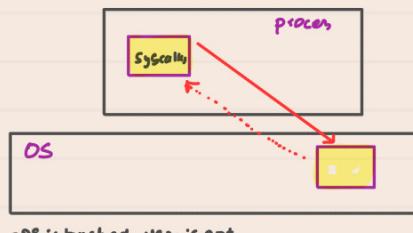
millions of lines of OS code

/ OS code



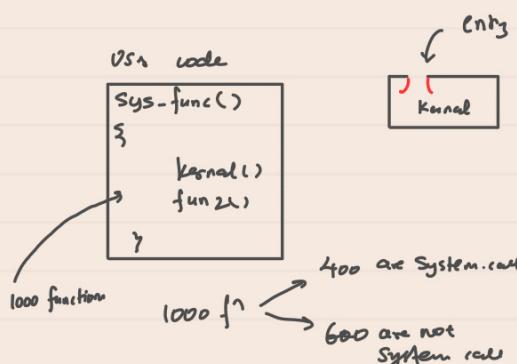
function call

Caller and callee are in the same process
- Same user
- Same "domain of trust"



- OS is trusted. User is not
- OS have super privilege. User not
- must take measures to prevent abuse

Some imp piece of code called Kernel
↓
Some critical fn that OS provide

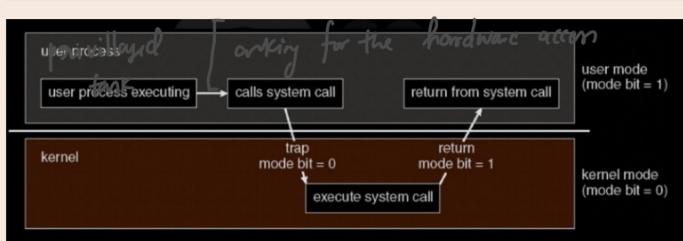


Entry point of Kernel

- System Calls
- interrupt
- exception

all the system calls are part of kernel

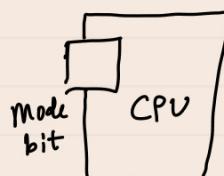
Q. how user ask for OS service?
using system call



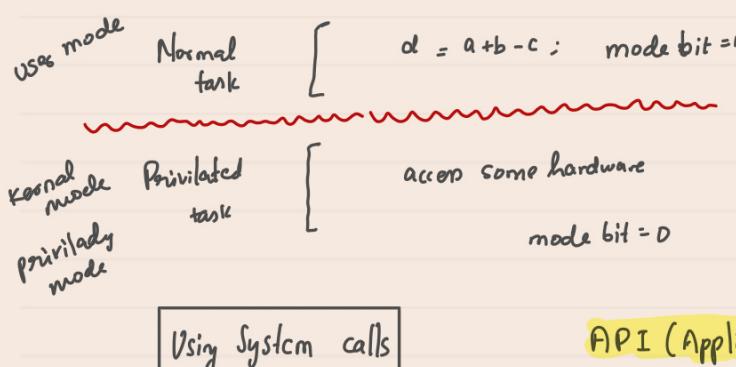
User Trust Issue

- Most of the users are not experts in computers and do not understand how hardware works. Thus they may not know how to properly manage the resources.
- An user (expert or not) can make incorrect operations while managing the resources (like handling h/w)
- An user can be malicious i.e. he/she can deliberately make incorrect operations. This can lead to system failure.

Solution from OS side: Never trust a user while dealing with critical tasks



Suppose user wants to access both



Using System calls

User Mode

- Compute operations like addition, Sub, etc
- Reading and writing from memory

Kernel mode

- CPU and memory management
- I/O handling

Processor Modes

- The OS must restrict what a user process can do
 - What instructions can execute
 - What portions of the address space are accessible
- Supervisor mode (or kernel mode)
 - Can execute any instructions in the instruction set
 - Including halting the processor, changing mode bit, initiating I/O
 - Can access any memory location in the system
 - Including code and data in the OS address space
- User mode
 - Restricted capabilities
 - Cannot execute privileged instructions
 - Cannot directly reference code or data in OS address space
 - Any such attempt results in a fatal protection fault
 - Instead, access OS code and data indirectly via system calls

API (Application programming interface)

User processes usually do not directly involve system calls. Instead they use something called an Application Programming Interface (API). It is a set of functions, variables and other software construction that allows to perform the most needed operations and also to cooperate with the operating system.

Q. System calls must be used to

- x (a) modify global variable
- x (b) call a user-written fn
- ✓ (c) write a file
- x (d) All of the above

lectures

Q. What is a system call?

Various possible answers

request to kernel to perform a service

- Open file
- Execute a new file
- Create a new process
- Send a message to another process

(Programmer Perspective)
"call a function"

- fd = open("myfile.txt");
- System call looks like any other function call

entry point providing controlled mechanism to execute kernel code

System call execution

- Each system call has unique numeric identifiers
- OS has a system call table that maps number to functionality required.

e.g.

High-level fork();
 ↓
 Compiler
 ↓
 Assembly mov rax, 57;
 ↓
 { syscall;
 ↓
 we have system call here

Rax 57
 ↓
 System call number

- o When invoking a system call, user places system call number and associated parameters in a specified register, then execute the syscall instruction.

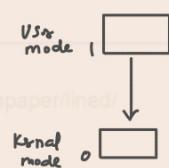
① Program calls Wrapper function in C Library

② Wrapper function

- o package syscall arguments into registers / stacks
- o puts (unique) syscall number into a register

③ Wrapper flips CPU to Kernel Mode (User-mode → Kernel-mode)

- o Execute special machine instruction (e.g. syscall)



- Main effect: CPU now can touch memory marked as accessible only in kernel mode

④ Kernel execute Syscall handler:

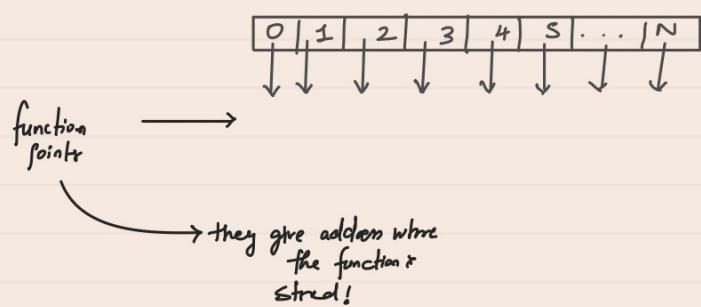
invokes service routine, corresponding to syscall number

- Do the real work generate result status

- Places return values from services routine in a register
- switch back to user mode, passing control back to wrapper
 - (Kernel mode → User-mode)

System Call table

- Syscall table is usually implemented as an array of function pointers, where each function implement one system call



$\text{foo}()$ $\xrightarrow{\quad}$ foo (an function pointer)
 \equiv

Q. Whenever we see "Syscall" instruction then what are the steps??

Q. True or false? System calls are as fast as function calls? Explain briefly??

False. System calls have to store syscall (unique) number and appropriate arguments then switch mode bit then branch through arrays of function points, then go to address execute and return value and switch back. takes hundreds of cycle. Whereas (assume) function call, just need to follow function point which is available in user mode address only.

Types of System call

CDA interpret Q's → Syscall

overlap!!

process control

- Create process, terminate process
- end, abort
- load, execute
- Get process attributes, set process attributes

file management

- create file, delete file
- open, close file
- read, write

Device management

- request device, release device

information maintenance

- get time or date, set time or date

Communication

production

Threads

Sometimes, a single application may required to perform similar tasks, in such cases we **don't** want to have several processes.

Why?

- * context switching
- * different address space

Solution?

Threads.

example

Word process

- spell checker
- write something



They share! unlike different processes they want everything for themselves!

for similar tasks, if we have different processes
(instead of threads) then

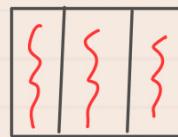
- IPC (interprocess communication)
is required

Types of thread process

- process context switching
is expensive than
thread context switching

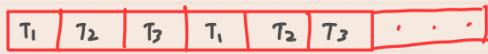


Single
threaded
process



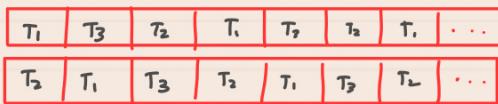
Multithreaded
processes

Concurrency:



"concurrent"

parallel
(need)
 $> 1 \text{ CPU}$



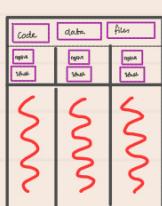
"parallel"

Q. Why do need to create single thread?

If could be you want just prefer one
similar task, 1 process ??
1 thread ??



Single
threaded
process!



Multithreaded
process!

{ shared!

} NOT shared OO!

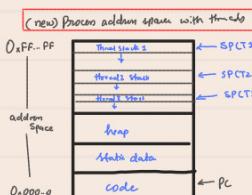
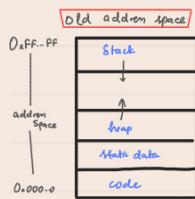
WAIT! WAIT! threads are like process they run individually in single CPU! So Now question?

What can you can't share??



What abt heap?

What abt it?
its space acquired
at runtime!



Types of thread

User level

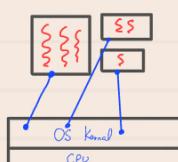
- When threads are created using function call

- have TCB (Thread Control Block)

Kernal Level

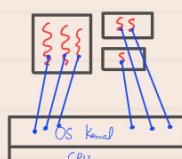
- When threads are created using System call

- have TCB (Thread Control Block)



Kernel knows only
three process.
Kernel don't know
how many threads
inside process the
are

OS will schedule the
process if doesn't know
about threads, then
effect



Kernal know about internal
threads.

OS could Schedule threads.
More efficient use of hardware.

Threads in
User level

CPU

Can we take help
from kernel scheduler
to schedule the threads??
⇒ No

> i(user) might be having my own
preference to schedule

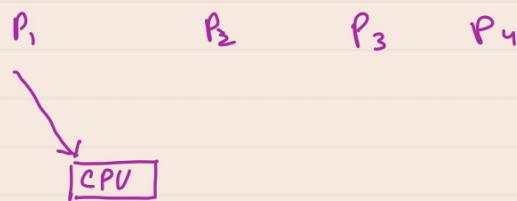
Q. I can not take help from OS scheduler then how am I supposed to schedule user level threads??

Ans: Thread Scheduler implemented in user space. (Repetitive task)
We have thread libraries, which has everything

Q. Can user level threads run kernel code?

Ans: Yes, we can definitely have systemcall

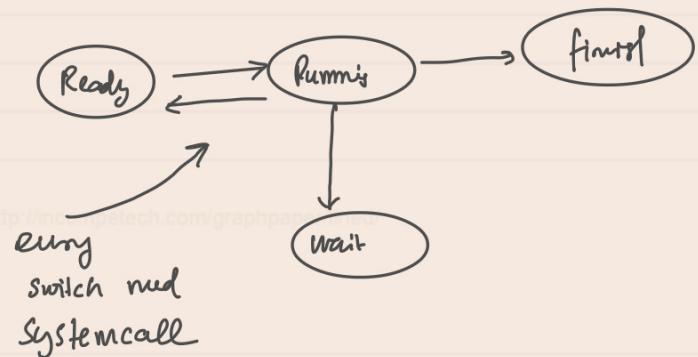
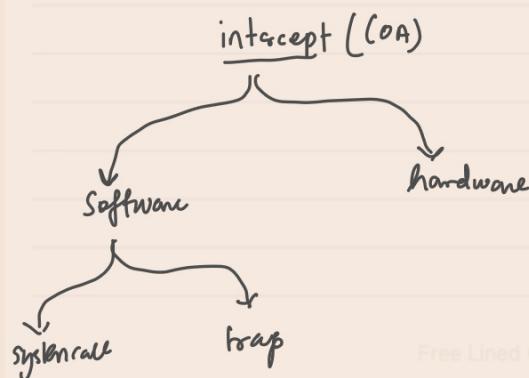
Q.



> P₁ is running
> We want to schedule other process now
> we will first run "scheduler"
How to Schedule, scheduler itself

There will be some interrupt

Sched; ← instruction



Q. What is one major advantage and disadvantage of User level threads?

① User flexibility

- fast to create (No systemcall needed)
- kernel doesn't have to maintain different threads. Hence context switching not

② efficiency could take hit

- if no of the threads is block then all the threads gets blocked, whole process gets blocked

needed

T/F

Q1) The OS provide the illusion to each thread that it has its own address space?

Ans: false, not completely.

Q2) With kernel threads, multiple threads from the same process can be scheduled on multiple CPUs simultaneously?

Ans: True

Q3) If a user level thread is blocked for I/O operation, then kernel of OS will perform context switching to run another user-level-thread which is not blocked??

Ans: False

2. Processes and threads. Suppose that three processes exist in a system, as described table below. Suppose that the system uses preemptive, round-robin scheduling, and that T_{11} is running when the quantum expires.

Process	Threads within the process
P_1	T_{11}, T_{12}, T_{13}
P_2	T_{21}, T_{22}
P_3	T_{31}

(a) (4) If the threads are implemented entirely at the user level (with no support from the operating system), which threads might possibly execute at the beginning of the next quantum?
 T_{21}, T_{22} or T_{31} . The OS is unaware of the threads within process P_1

(b) (4) If threads are supported by the operating system (i.e., lightweight processes), which threads might possibly execute at the beginning of the next quantum?
any except T_{11}

(c) (2) How does your answer change if the system has two processors?
It actually doesn't change. In the first case, the kernel will not know about the user's threads, so no other choices can be made. In the second case, the kernel will have full knowledge of the user's threads and will make the same decision (of course, realizing that the OS should not run the same thread on simultaneously on multiple processors).

10:25 am ✓

Ⓐ T_{21}, T_{22}, T_{31}

Ⓑ any except T_{11}

Ⓒ if all user level
then
 T_{31} will run
and
either T_{21} or T_{22}

if kernel level
except T_{11}
any other two
threads will run

- Q. A blocking kernel-scheduled thread blocks all the threads in the process - False
- Q. Threads are cheaper to context switch than processes - True
- Q. A blocking user-level thread blocks the process - True

Indicate which statements are True, and which are False, regarding a single process and its threads:

38. In an operating system environment that operates with User-Level threads only (no Kernel-Level threads), the various threads associated with a single process all share a common context, including a common PC.
False
39. In an operating system environment that operates with Kernel-Level threads only (no User-Level threads), the various threads associated with a single process all share a common context, including a common PC.
False
40. The different threads associated with a single process all share a single processor context, i.e., the value of the Processor Status Word is the same for all such threads.
??
41. Thread creation takes more time than process creation.
false
42. Each thread associated with a single process has exclusive ownership of any files that it has opened; they are not available to other threads associated with the same process.
false
43. Each thread associated with a single process has a separate text (program code) area.
10:34 am ✓
false

HOMEWORK (GATE Questions)

read question correctly!

- Q. (4 marks)
Suppose that two long running processes, P_1 and P_2 , are running in a system. Neither program performs any system calls that might cause it to block, and there are no other processes in the system. P_1 has 2 threads and P_2 has 1 thread.

(i) What percentage of CPU time will P_1 get if the threads are *kernel threads*? Explain your answer.

(ii) What percentage of CPU time will P_1 get if the threads are *user threads*? Explain your answer.

Sample Answer:

(i) If the threads are kernel threads, they are independently scheduled and each of the three threads will get a share of the CPU. Thus, the 2 threads of P_1 will get $2/3$ of the CPU time. That is, P_1 will get 66% of the CPU.

(ii) If the threads are user threads, the threads of each process map to one kernel thread, so each process will get a share of the CPU. The kernel is unaware that P_1 has two threads. Thus, P_1 will get 50% of the CPU.

5:17 am ✓

Trying

P_1 P_2
{} {}

(No syscalls that would block)
(No other processes)

(i) If CPU, assume round robin
Then $\frac{1}{2}$ cause CPU knows that

$$(II) \frac{1}{2} \cdot \left(\frac{1}{2} \right)$$

in 1 process P_1 has
that changes

1 process

$$\frac{2}{3}$$

process 1 has 2 threads
and it can schedule them

$$\frac{1}{4}$$

$$\frac{1}{2}$$

GATE CSE 2004

44 Consider the following statements with respect to user-level threads and kernel-supported threads

- I. context switch is faster with kernel-supported threads
- II. for user-level threads, a system call can block the entire process
- III. Kernel supported threads can be scheduled independently
- IV. User level threads are transparent to the kernel

Which of the above statements are true?

A. (II), (III) and (IV) only
 B. (II) and (III) only
 C. (I) and (III) only
 D. (I) and (II) only

5:22 am ✓

(I) F

(II) T

(III) T

(IV) T

Ans: (A)

S

transparent: invisible

GATE 2007

32 Consider the following statements about user level threads and kernel level threads. Which one of the following statements is FALSE?

- A. Context switch time is longer for kernel level threads than for user level threads.
- B. User level threads do not need any hardware support.
- C. Related kernel level threads can be scheduled on different processors in a multi-processor system.
- D. Blocking one kernel level thread blocks all related threads.

A. T

B. F, User level thread can call Systemcall as needed

C. T

D. F, in the case with user level threads

Ans: (B),(D)

B - True, User level threads don't need any hardware support.

Kernel level thread - GOD - implemented and recognized by OS

Requires Hardware Support

Independent

More context switching time (b/c more difficult hard implementation)

User level thread - HUMAN

USER LEVEL THREAD	KERNEL LEVEL THREAD
User thread are implemented by users.	kernel threads are implemented by OS.
OS doesn't recognize user level threads.	Kernel threads are recognized by OS.
Implementation of User threads is easy.	Implementation of Kernel threads is complicated.
Context switch time is less.	Context switch time is more.
Context switch requires no hardware support.	Hardware support is needed.
If one user level thread performs blocking operation then entire process will be blocked.	If one kernel thread performs blocking operation then another thread can continue execution.
Example : Java thread, POSIX threads.	Example : Windows Solaris.

5:37 am ✓

GATE CSE 2011

A thread is usually defined as a light weight process because an Operating System (OS) maintains smaller data structures for a thread than for a process. In relation to this, which of the following statements is correct?

- A. OS maintains only scheduling and accounting information for each thread
- B. OS maintains only CPU registers for each thread
- C. OS does not maintain virtual memory state for each thread
- D. OS does not maintain a separate stack for each thread

A. Thread (stack, reg) {heap, data, code}

B. False,

C. TRUE,  inside a process there is thread so maybe

D. False

OS, on per thread basis, maintain only two things  CPU registers
Stack Space

Even TLB and page tables are also shared since they belong to same process
 Scheduling help Accounting
(unable to understand)

GATE 2014

Which one of the following is FALSE?

- A. User level threads are not scheduled by the kernel.
- B. When a user level thread is blocked, all other threads of its process are blocked.
- C. Context switching between user level threads is faster than context switching between kernel level threads.
- D. Kernel level threads cannot share the code segment.

A. True, kernel level threads are scheduled by kernel

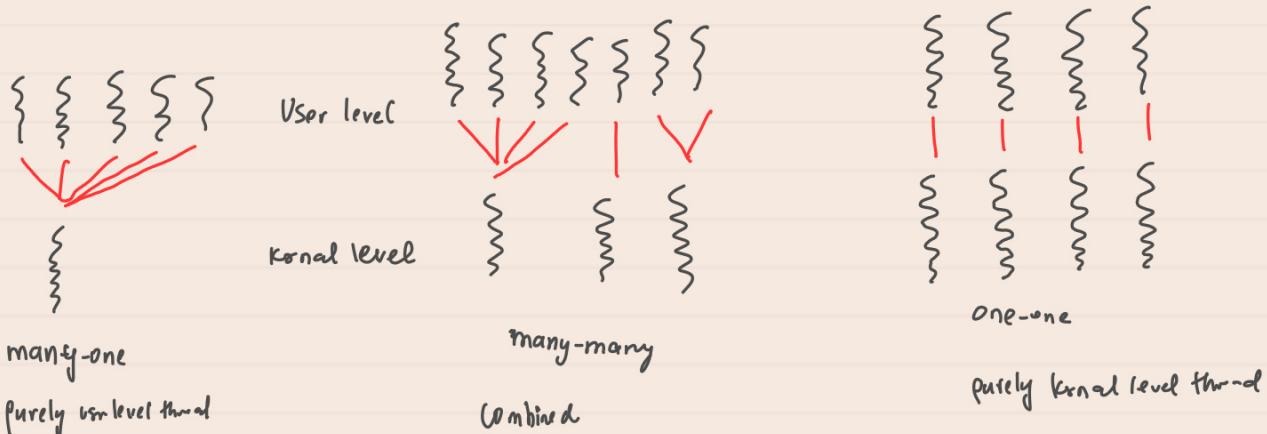
B. True

C. True

D. False, be it user level or kernel level, threads share code, heap, data with common thread of self process.

Multithreading Models

- Many-one
- one to one
- many to many



- Q.  Suppose: we want to run V3 even if V1 is blocked then

④ we will map u_1, u_3 to one kernel level thread
 ⑤ we will map v_1, v_3 to different kernel level thread

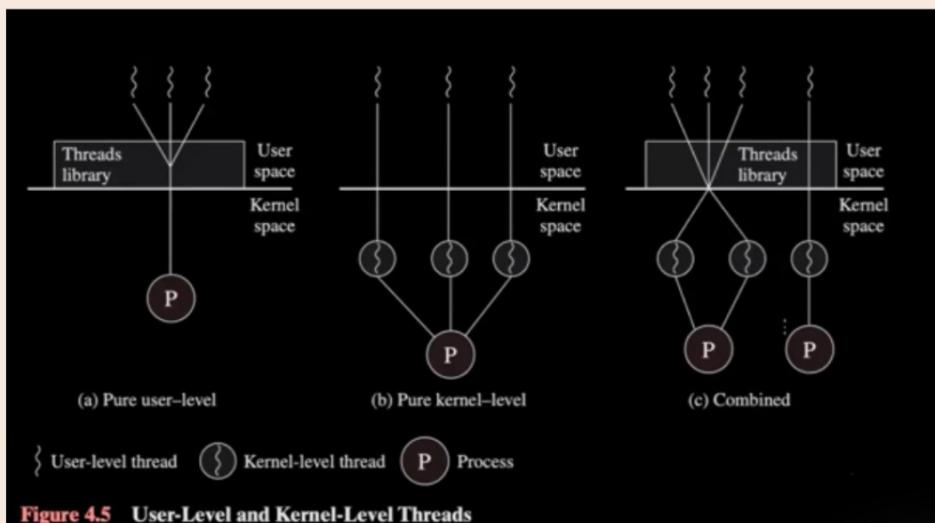


Figure 4.5 User-Level and Kernel-Level Threads

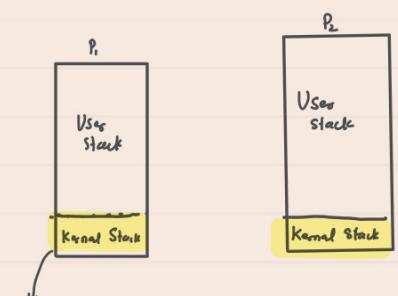
CONTENT SWITCHING

due to any reason!



Content Switching

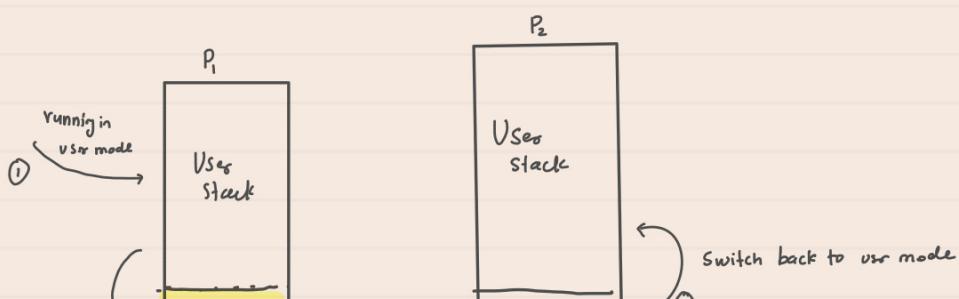
- Saves state of a process before switching to another process
- Restore original process state when switching back

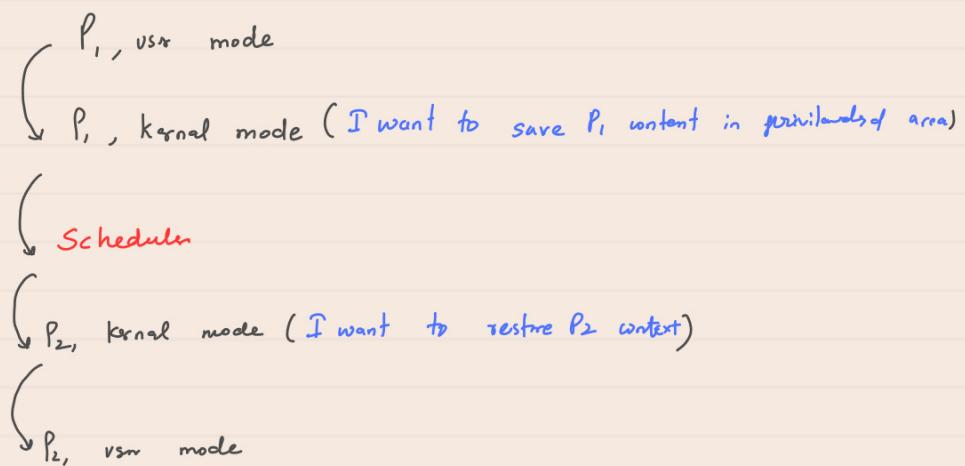
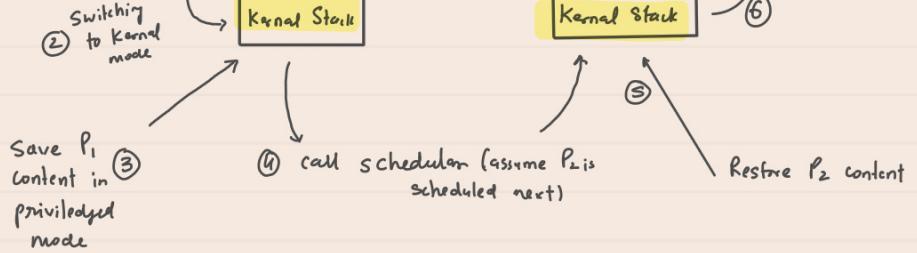


We can push the register value here at the time of content switching

Why are we using Stack data structure?

Cause stack isn't fix, it can grow and shorten itself realtime depending upon the required unlike fixed array size allocated for storing reg. values.





Free Lined Graph Paper from <http://incompetech.com/graphpaperlined/>

Q. Context switching only happens when process is already in kernel mode

Ans: True, you can't directly switch from P₁(user mode) to P₂

(Thread quantum)